

R version 4.4.2 (2024-10-31 ucrt) -- "Pile of Leaves"
 Copyright (C) 2024 The R Foundation for Statistical Computing
 Platform: x86_64-w64-mingw32/x64

R is free software and comes with ABSOLUTELY NO WARRANTY.
 You are welcome to redistribute it under certain conditions.
 Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
 Type 'contributors()' for more information and
 'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
 'help.start()' for an HTML browser interface to help.
 Type 'q()' to quit R.

[Previously saved workspace restored]

```
> library(rpart)
Warning message:
package 'rpart' was built under R version 4.4.3
> dt_model <- rpart(Species ~ ., data = iris_train, method = "class")
> summary(dt_model)
Call:
rpart(formula = Species ~ ., data = iris_train, method = "class")
n= 105
```

	CP	nsplit	rel error	xerror	xstd
1	0.4846154	0	1.00000000	1.00000000	0.07655590
2	0.0100000	2	0.03076923	0.04615385	0.02626351

Variable importance

PetalWidth	PetalLength	SepalLength	SepalWidth
34	32	21	13

Node number 1: 105 observations, complexity param=0.4846154
 predicted class=virginica expected loss=0.6190476 P(node) =1
 class counts: 32 33 40
 probabilities: 0.305 0.314 0.381
 left son=2 (67 obs) right son=3 (38 obs)
 Primary splits:
 PetalWidth < 1.75 to the left, improve=34.23511, (0 missing)
 PetalLength < 2.45 to the left, improve=33.47371, (0 missing)
 SepalLength < 5.5 to the left, improve=21.16035, (0 missing)
 SepalWidth < 3.05 to the right, improve=13.84740, (0 missing)
 Surrogate splits:
 PetalLength < 4.75 to the left, agree=0.952, adj=0.868, (0 split)
 SepalLength < 6.15 to the left, agree=0.838, adj=0.553, (0 split)

Node number 2: 67 observations, complexity param=0.4846154
 predicted class=versicolor expected loss=0.5074627 P(node) =0.6380952
 class counts: 32 33 2
 probabilities: 0.478 0.493 0.030
 left son=4 (32 obs) right son=5 (35 obs)
 Primary splits:
 PetalLength < 2.45 to the left, improve=31.63156, (0 missing)
 PetalWidth < 0.8 to the left, improve=31.63156, (0 missing)
 SepalWidth < 3.05 to the right, improve=21.47441, (0 missing)
 SepalLength < 5.5 to the left, improve=15.35073, (0 missing)
 Surrogate splits:
 PetalWidth < 0.8 to the left, agree=1.000, adj=1.000, (0 split)
 SepalWidth < 3.05 to the right, agree=0.910, adj=0.812, (0 split)
 SepalLength < 5.5 to the left, agree=0.836, adj=0.656, (0 split)

Node number 3: 38 observations
 predicted class=virginica expected loss=0 P(node) =0.3619048
 class counts: 0 0 38

```
probabilities: 0.000 0.000 1.000
```

```
Node number 4: 32 observations
```

```
predicted class=setosa      expected loss=0   P(node) =0.3047619
```

```
class counts:    32      0      0
```

```
probabilities: 1.000 0.000 0.000
```

```
Node number 5: 35 observations
```

```
predicted class=versicolor expected loss=0.05714286 P(node) =0.3333333
```

```
class counts:    0    33    2
```

```
probabilities: 0.000 0.943 0.057
```

```
> # Predict classes for the training data
```

```
> train_pred <- predict(dt_model, iris_train, type = "class")
```

```
>
```

```
> # Build confusion matrix for training data
```

```
> train_conf_matrix <- table(Predicted = train_pred, Actual = iris_train$Species)
```

```
>
```

```
> # View confusion matrix
```

```
> train_conf_matrix
```

```
      Actual
```

```
Predicted   setosa versicolor virginica
```

```
setosa       32         0         0
```

```
versicolor   0         33         2
```

```
virginica    0         0        38
```

```
> # Predict classes for the testing data
```

```
> test_pred <- predict(dt_model, iris_test, type = "class")
```

```
>
```

```
> # Build confusion matrix for testing data
```

```
> test_conf_matrix <- table(Predicted = test_pred, Actual = iris_test$Species)
```

```
>
```

```
> # View confusion matrix
```

```
> test_conf_matrix
```

```
      Actual
```

```
Predicted   setosa versicolor virginica
```

```
setosa       12         0         0
```

```
versicolor   0        21         2
```

```
virginica    0         0        10
```

```
> train_accuracy <- sum(diag(train_conf_matrix)) / sum(train_conf_matrix)
```

```
> train_accuracy
```

```
[1] 0.9809524
```

```
> test_accuracy <- sum(diag(test_conf_matrix)) / sum(test_conf_matrix)
```

```
> test_accuracy
```

```
[1] 0.9555556
```

```
> # Sensitivity for training set
```

```
> train_sensitivity <- diag(train_conf_matrix) / colSums(train_conf_matrix)
```

```
> train_sensitivity
```

```
      setosa versicolor virginica
```

```
1.00    1.00    0.95
```

```
> # Sensitivity for testing set
```

```
> test_sensitivity <- diag(test_conf_matrix) / colSums(test_conf_matrix)
```

```
> test_sensitivity
```

```
      setosa versicolor virginica
```

```
1.0000000 1.0000000 0.8333333
```

```
> # Function to calculate specificity
```

```
> calc_specificity <- function(conf_matrix) {
```

```
+   spec <- numeric(nrow(conf_matrix))
```

```
+   for (i in 1:nrow(conf_matrix)) {
```

```
+     true_negatives <- sum(conf_matrix[-i, -i])
```

```
+     false_positives <- sum(conf_matrix[i, -i])
```

```
+     spec[i] <- true_negatives / (true_negatives + false_positives)
```

```
+   }
```

```
+   return(spec)
```

```
+ }
```

```
>
```

```
> # Specificity for training set
```

```
> train_specificity <- calc_specificity(train_conf_matrix)
```

```
> train_specificity
```

```
[1] 1.0000000 0.9722222 1.0000000
```

```
>
```

```

> # Specificity for testing set
> test_specificity <- calc_specificity(test_conf_matrix)
> test_specificity
[1] 1.0000000 0.9166667 1.0000000
> install.packages("caret")
Installing package into 'C:/Users/Simon/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)
--- Please select a CRAN mirror for use in this session ---
trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/caret_7.0-1.zip'
Content type 'application/zip' length 3603291 bytes (3.4 MB)
downloaded 3.4 MB

package 'caret' successfully unpacked and MD5 sums checked
Warning: cannot remove prior installation of package 'caret'
Warning: restored 'caret'

The downloaded binary packages are in
      C:\Users\Simon\AppData\Local\Temp\RtmpQNRidw\downloaded_packages
Warning message:
In file.copy(savedcopy, lib, recursive = TRUE) :
  problem copying C:\Users\Simon\AppData\Local\R\win-library\4.4\00LOCK\caret\libs\x64\caret.dll
to C:\Users\Simon\AppData\Local\R\win-library\4.4\caret\libs\x64\caret.dll: Permission denied
> library(caret)
Loading required package: ggplot2
Loading required package: lattice
Warning message:
package 'caret' was built under R version 4.4.3
> # Create a vector to store accuracies
> accuracies <- c()
>
> # List of k values
> k_values <- c(5, 10, 15, 20)
>
> # Loop through each k
> for (k in k_values) {
+   # Set up k-fold cross validation
+   cv_control <- trainControl(method = "cv", number = k)
+   # Train the Decision Tree model using caret's train() function
+   cv_model <- train(Species ~ .,
+                     data = iris_o_sample,
+                     method = "rpart",
+                     trControl = cv_control)
+   # Store the mean accuracy
+   accuracies <- c(accuracies, max(cv_model$results$Accuracy))
+ }
>
> # View the results
> data.frame(K = k_values, Accuracy = accuracies)
  K Accuracy
1  5 0.9666370
2 10 0.9665476
3 15 0.9715320
4 20 0.9648810
> # Simple plot
> plot(k_values, accuracies, type = "b", pch = 19,
+       xlab = "K (Number of Folds)",
+       ylab = "Accuracy",
+       main = "Accuracy vs K-Folds for Decision Tree")
> install.packages("e1071")
Installing package into 'C:/Users/Simon/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)
trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/e1071_1.7-16.zip'
Content type 'application/zip' length 674453 bytes (658 KB)
downloaded 658 KB

package 'e1071' successfully unpacked and MD5 sums checked
Warning: cannot remove prior installation of package 'e1071'

```

Warning: restored 'e1071'

The downloaded binary packages are in

C:\Users\Simon\AppData\Local\Temp\RtmpQNRidw\downloaded_packages

Warning message:

In file.copy(savedcopy, lib, recursive = TRUE) :

problem copying C:\Users\Simon\AppData\Local\R\win-library\4.4\00LOCK\e1071\libs\x64\e1071.dll
to C:\Users\Simon\AppData\Local\R\win-library\4.4\e1071\libs\x64\e1071.dll: Permission denied

> library(e1071)

Warning message:

package 'e1071' was built under R version 4.4.3

> # Train the Naive Bayes model

> nb_model <- naiveBayes(Species ~ ., data = iris_train)

>

> # Predict on training set

> train_pred_nb <- predict(nb_model, iris_train)

>

> # Predict on testing set

> test_pred_nb <- predict(nb_model, iris_test)

>

> # Build confusion matrices

> train_conf_matrix_nb <- table(Predicted = train_pred_nb, Actual = iris_train\$Species)

> test_conf_matrix_nb <- table(Predicted = test_pred_nb, Actual = iris_test\$Species)

>

> # View them

> train_conf_matrix_nb

	Actual		
Predicted	setosa	versicolor	virginica
setosa	32	0	0
versicolor	0	32	2
virginica	0	1	38

> test_conf_matrix_nb

	Actual		
Predicted	setosa	versicolor	virginica
setosa	12	0	0
versicolor	0	19	0
virginica	0	2	12

> train_accuracy_nb <- sum(diag(train_conf_matrix_nb)) / sum(train_conf_matrix_nb)

> train_accuracy_nb

[1] 0.9714286

> test_accuracy_nb <- sum(diag(test_conf_matrix_nb)) / sum(test_conf_matrix_nb)

> test_accuracy_nb

[1] 0.9555556

> # Function to calculate precision, recall, F1 for each class

> calc_metrics <- function(conf_matrix) {

+ precision <- diag(conf_matrix) / rowSums(conf_matrix)

+ recall <- diag(conf_matrix) / colSums(conf_matrix)

+ f1 <- 2 * precision * recall / (precision + recall)

+ data.frame(Precision = precision, Recall = recall, F1 = f1)

+ }

>

> # Metrics for training set

> train_metrics_nb <- calc_metrics(train_conf_matrix_nb)

> train_metrics_nb

	Precision	Recall	F1
setosa	1.0000000	1.000000	1.0000000
versicolor	0.9411765	0.969697	0.9552239
virginica	0.9743590	0.950000	0.9620253

>

> # Metrics for testing set

> test_metrics_nb <- calc_metrics(test_conf_matrix_nb)

> test_metrics_nb

	Precision	Recall	F1
setosa	1.0000000	1.0000000	1.0000000
versicolor	1.0000000	0.9047619	0.9500000
virginica	0.8571429	1.0000000	0.9230769

> install.packages("class")

Installing package into 'C:/Users/Simon/AppData/Local/R/win-library/4.4'

(as 'lib' is unspecified)

trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/class_7.3-23.zip'

Content type 'application/zip' length 99782 bytes (97 KB)
downloaded 97 KB

package 'class' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Simon\AppData\Local\Temp\RtmpQNRidw\downloaded_packages

```
> library(class)
```

```
Error in value[[3L]](cond) :
```

```
Package 'class' version 7.3.22 cannot be unloaded:
```

```
Error in unloadNamespace(package) : namespace 'class' is imported by 'ipred', 'e1071' so cannot be unloaded
```

```
In addition: Warning message:
```

```
package 'class' was built under R version 4.4.3
```

```
> # Set k = 5 neighbors (you can adjust later if needed)
```

```
> k_value <- 5
```

```
>
```

```
> # Train and predict on training set
```

```
> train_pred_knn <- knn(train = iris_train[, -5],
```

```
+ test = iris_train[, -5],
```

```
+ cl = iris_train$Species,
```

```
+ k = k_value)
```

```
Error in knn(train = iris_train[, -5], test = iris_train[, -5], cl = iris_train$Species, :  
could not find function "knn"
```

```
>
```

```
> # Train and predict on testing set
```

```
> test_pred_knn <- knn(train = iris_train[, -5],
```

```
+ test = iris_test[, -5],
```

```
+ cl = iris_train$Species,
```

```
+ k = k_value)
```

```
Error in knn(train = iris_train[, -5], test = iris_test[, -5], cl = iris_train$Species, :  
could not find function "knn"
```

```
>
```

```
> # Build confusion matrices
```

```
> train_conf_matrix_knn <- table(Predicted = train_pred_knn, Actual = iris_train$Species)
```

```
Error: object 'train_pred_knn' not found
```

```
> test_conf_matrix_knn <- table(Predicted = test_pred_knn, Actual = iris_test$Species)
```

```
Error: object 'test_pred_knn' not found
```

```
>
```

```
> # View the matrices
```

```
> train_conf_matrix_knn
```

```
Error: object 'train_conf_matrix_knn' not found
```

```
> test_conf_matrix_knn
```

```
Error: object 'test_conf_matrix_knn' not found
```

```
> train_pred_knn <- class::knn(train = iris_train[, -5],
```

```
+ test = iris_train[, -5],
```

```
+ cl = iris_train$Species,
```

```
+ k = k_value)
```

```
>
```

```
> test_pred_knn <- class::knn(train = iris_train[, -5],
```

```
+ test = iris_test[, -5],
```

```
+ cl = iris_train$Species,
```

```
+ k = k_value)
```

```
> # Build confusion matrices
```

```
> train_conf_matrix_knn <- table(Predicted = train_pred_knn, Actual = iris_train$Species)
```

```
> test_conf_matrix_knn <- table(Predicted = test_pred_knn, Actual = iris_test$Species)
```

```
>
```

```
> # View them
```

```
> train_conf_matrix_knn
```

```
Actual
```

```
Predicted setosa versicolor virginica
```

```
setosa 32 0 0
```

```
versicolor 0 32 2
```

```
virginica 0 1 38
```

```
> test_conf_matrix_knn
```

```
Actual
```

```
Predicted setosa versicolor virginica
```

```
setosa 12 0 0
```

```
versicolor 0 18 0
```

```
virginica 0 3 12
```

```

> train_accuracy_knn <- sum(diag(train_conf_matrix_knn)) / sum(train_conf_matrix_knn)
> train_accuracy_knn
[1] 0.9714286
> test_accuracy_knn <- sum(diag(test_conf_matrix_knn)) / sum(test_conf_matrix_knn)
> test_accuracy_knn
[1] 0.9333333
> # Function to calculate precision, recall, F1 for each class
> calc_metrics <- function(conf_matrix) {
+   precision <- diag(conf_matrix) / rowSums(conf_matrix)
+   recall <- diag(conf_matrix) / colSums(conf_matrix)
+   f1 <- 2 * precision * recall / (precision + recall)
+   data.frame(Precision = precision, Recall = recall, F1 = f1)
+ }
>
> # Metrics for training set
> train_metrics_knn <- calc_metrics(train_conf_matrix_knn)
> train_metrics_knn
      Precision    Recall      F1
setosa      1.0000000 1.000000 1.0000000
versicolor 0.9411765 0.969697 0.9552239
virginica   0.9743590 0.950000 0.9620253
>
> # Metrics for testing set
> test_metrics_knn <- calc_metrics(test_conf_matrix_knn)
> test_metrics_knn
      Precision    Recall      F1
setosa      1.0 1.0000000 1.0000000
versicolor  1.0 0.8571429 0.9230769
virginica    0.8 1.0000000 0.8888889
> install.packages("RWeka")
Installing package into 'C:/Users/Simon/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)
also installing the dependencies 'RWekajars', 'rJava'

trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/RWekajars_3.9.3-2.zip'
Content type 'application/zip' length 10032953 bytes (9.6 MB)
downloaded 9.6 MB

trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/rJava_1.0-11.zip'
Content type 'application/zip' length 835913 bytes (816 KB)
downloaded 816 KB

trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/RWeka_0.4-46.zip'
Content type 'application/zip' length 541736 bytes (529 KB)
downloaded 529 KB

package 'RWekajars' successfully unpacked and MD5 sums checked
package 'rJava' successfully unpacked and MD5 sums checked
package 'RWeka' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
      C:\Users\Simon\AppData\Local\Temp\RtmpQNRidw\downloaded_packages
> library(RWeka)
Warning message:
package 'RWeka' was built under R version 4.4.3
> # Train JRip model
> rule_model <- JRip(Species ~ ., data = iris_train)
Error in .jcall(o, "Ljava/lang/Class;", "getClass") :
  weka.core.UnsupportedAttributeTypeException: weka.classifiers.rules.JRip: Cannot handle string
class!
>
> # Predict on training set
> train_pred_rule <- predict(rule_model, iris_train)
Error: object 'rule_model' not found
>
> # Predict on testing set
> test_pred_rule <- predict(rule_model, iris_test)
Error: object 'rule_model' not found
>
> # Build confusion matrices

```

```

> train_conf_matrix_rule <- table(Predicted = train_pred_rule, Actual = iris_train$Species)
Error: object 'train_pred_rule' not found
> test_conf_matrix_rule <- table(Predicted = test_pred_rule, Actual = iris_test$Species)
Error: object 'test_pred_rule' not found
>
> # View matrices
> train_conf_matrix_rule
Error: object 'train_conf_matrix_rule' not found
> test_conf_matrix_rule
Error: object 'test_conf_matrix_rule' not found
> # Make sure Species is a factor
> iris_train$Species <- as.factor(iris_train$Species)
> iris_test$Species <- as.factor(iris_test$Species)
> # Train JRip model
> rule_model <- JRip(Species ~ ., data = iris_train)
>
> # Predict on training set
> train_pred_rule <- predict(rule_model, iris_train)
>
> # Predict on testing set
> test_pred_rule <- predict(rule_model, iris_test)
>
> # Build confusion matrices
> train_conf_matrix_rule <- table(Predicted = train_pred_rule, Actual = iris_train$Species)
> test_conf_matrix_rule <- table(Predicted = test_pred_rule, Actual = iris_test$Species)
>
> # View matrices
> train_conf_matrix_rule
      Actual
Predicted setosa versicolor virginica
setosa      32           0           0
versicolor   0          32           1
virginica    0           1          39
> test_conf_matrix_rule
      Actual
Predicted setosa versicolor virginica
setosa      12           0           0
versicolor   0          21           2
virginica    0           0          10
> train_accuracy_rule <- sum(diag(train_conf_matrix_rule)) / sum(train_conf_matrix_rule)
> train_accuracy_rule
[1] 0.9809524
> test_accuracy_rule <- sum(diag(test_conf_matrix_rule)) / sum(test_conf_matrix_rule)
> test_accuracy_rule
[1] 0.9555556
> # Function to calculate precision, recall, F1
> calc_metrics <- function(conf_matrix) {
+   precision <- diag(conf_matrix) / rowSums(conf_matrix)
+   recall <- diag(conf_matrix) / colSums(conf_matrix)
+   f1 <- 2 * precision * recall / (precision + recall)
+   data.frame(Precision = precision, Recall = recall, F1 = f1)
+ }
>
> # Metrics for training set
> train_metrics_rule <- calc_metrics(train_conf_matrix_rule)
> train_metrics_rule
      Precision   Recall      F1
setosa      1.000000 1.000000 1.000000
versicolor  0.969697 0.969697 0.969697
virginica   0.975000 0.975000 0.975000
>
> # Metrics for testing set
> test_metrics_rule <- calc_metrics(test_conf_matrix_rule)
> test_metrics_rule
      Precision   Recall      F1
setosa      1.000000 1.000000 1.000000
versicolor  0.9130435 1.000000 0.9545455
virginica   1.000000 0.8333333 0.9090909
> # Model names
> model_names <- c("Decision Tree", "Naive Bayes", "kNN", "Rule-Based")

```

```

>
> # Corresponding testing accuracies (fill in your real values)
> accuracies <- c(0.9556, 0.9556, 0.9333, 0.9556)
>
> # Make a data frame for plotting
> accuracy_df <- data.frame(Model = model_names, Accuracy = accuracies)
> # Bar plot of testing accuracies
> barplot(accuracy_df$Accuracy,
+         names.arg = accuracy_df$Model,
+         col = "skyblue",
+         ylim = c(0, 1),
+         main = "Comparison of Model Accuracies",
+         ylab = "Accuracy",
+         xlab = "Models")
> # Create a fake expanded dataset
> expanded_accuracy_df <- data.frame(
+   Model = rep(model_names, each = 5), # Fake 5 observations per model
+   Accuracy = c(
+     rep(0.9556, 5), # Decision Tree
+     rep(0.9556, 5), # Naive Bayes
+     rep(0.9333, 5), # kNN
+     rep(0.9556, 5) # Rule-Based
+   )
+ )
>
> # Real box plot
> boxplot(Accuracy ~ Model,
+         data = expanded_accuracy_df,
+         col = "lightgreen",
+         ylim = c(0, 1),
+         main = "Comparison of Model Accuracies (Box Plot Style)",
+         ylab = "Accuracy",
+         xlab = "Models")
> install.packages("pROC")
Installing package into 'C:/Users/Simon/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)
trying URL 'https://lib.stat.cmu.edu/R/CRAN/bin/windows/contrib/4.4/pROC_1.18.5.zip'
Content type 'application/zip' length 1168382 bytes (1.1 MB)
downloaded 1.1 MB

```

```

package 'pROC' successfully unpacked and MD5 sums checked
Warning: cannot remove prior installation of package 'pROC'
Warning: restored 'pROC'

```

```

The downloaded binary packages are in
  C:\Users\Simon\AppData\Local\Temp\RtmpQNRidw\downloaded_packages
Warning message:
In file.copy(savedcopy, lib, recursive = TRUE) :
  problem copying C:\Users\Simon\AppData\Local\R\win-library\4.4\00LOCK\pROC\libs\x64\pROC.dll to
  C:\Users\Simon\AppData\Local\R\win-library\4.4\pROC\libs\x64\pROC.dll: Permission denied
> install.packages("caret")
Warning: package 'caret' is in use and will not be installed
> library(pROC)
Type 'citation("pROC")' for a citation.

```

```

Attaching package: 'pROC'

```

```

The following objects are masked from 'package:stats':

```

```

  cov, smooth, var

```

```

Warning message:
package 'pROC' was built under R version 4.4.3
> library(caret)
> # DT Probabilities on testing set
> dt_probs <- predict(dt_model, iris_test, type = "prob")
> # NB Probabilities on testing set
> nb_probs <- predict(nb_model, iris_test, type = "raw")
> # Simulate JRip prediction probabilities
> rule_probs <- predict(rule_model, iris_test, type = "class")

```



```

> rule_probs <- model.matrix(~ rule_probs - 1) # Trick: create dummy matrix
> colnames(rule_probs) <- levels(iris_test$Species)
> # Create ROC curves
> dt_roc <- roc(iris_test$Species, dt_probs[, "setosa"], levels = rev(levels(iris_test$Species)))
Error in roc.default(iris_test$Species, dt_probs[, "setosa"], levels = rev(levels(iris_test$Species))) :
  'levels' argument must have length 2
> nb_roc <- roc(iris_test$Species, nb_probs[, "setosa"], levels = rev(levels(iris_test$Species)))
Error in roc.default(iris_test$Species, nb_probs[, "setosa"], levels = rev(levels(iris_test$Species))) :
  'levels' argument must have length 2
> rule_roc <- roc(iris_test$Species, rule_probs[, "setosa"], levels = rev(levels(iris_test$Species)))
Error in roc.default(iris_test$Species, rule_probs[, "setosa"], levels = rev(levels(iris_test$Species))) :
  'levels' argument must have length 2
>
> # Plot them together
> plot(dt_roc, col = "blue", lwd = 2, main = "ROC Curve - Setosa vs Others")
Error: object 'dt_roc' not found
> plot(nb_roc, col = "green", lwd = 2, add = TRUE)
Error: object 'nb_roc' not found
> plot(rule_roc, col = "red", lwd = 2, add = TRUE)
Error: object 'rule_roc' not found
>
> legend("bottomright", legend = c("Decision Tree", "Naive Bayes", "Rule-Based"),
+       col = c("blue", "green", "red"), lwd = 2)
> # Create binary labels: Setosa vs Not Setosa
> iris_test_binary <- ifelse(iris_test$Species == "setosa", "setosa", "not_setosa")
> iris_test_binary <- factor(iris_test_binary, levels = c("not_setosa", "setosa"))
> dt_roc <- roc(iris_test_binary, dt_probs[, "setosa"])
Setting levels: control = not_setosa, case = setosa
Setting direction: controls < cases
> nb_roc <- roc(iris_test_binary, nb_probs[, "setosa"])
Setting levels: control = not_setosa, case = setosa
Setting direction: controls < cases
> rule_roc <- roc(iris_test_binary, rule_probs[, "setosa"])
Setting levels: control = not_setosa, case = setosa
Setting direction: controls < cases
> # Plot DT ROC first
> plot(dt_roc, col = "blue", lwd = 2, main = "ROC Curve - Setosa vs Others")
>
> # Add NB ROC
> plot(nb_roc, col = "green", lwd = 2, add = TRUE)
>
> # Add Rule-Based ROC
> plot(rule_roc, col = "red", lwd = 2, add = TRUE)
>
> # Add a legend
> legend("bottomright", legend = c("Decision Tree", "Naive Bayes", "Rule-Based"),
+       col = c("blue", "green", "red"), lwd = 2)
> x<-read.table("/Users/Simon/Downloads/seeds_original.csv", header = T, sep=",")
> seeds_original<-data.frame(x)
> head(seeds_original)
  Area Perimeter Compactness LengthKernel WidthKernel AsymmetryCoefficient LengthKernelGroove Cl
ass
1 15.26      14.84      0.8710      5.763      3.312      2.221      5.220 K
ama
2 14.88      14.57      0.8811      5.554      3.333      1.018      4.956 K
ama
3 14.29      14.09      0.9050      5.291      3.337      2.699      4.825 K
ama
4 13.84      13.94      0.8955      5.324      3.379      2.259      4.805 K
ama
5 16.14      14.99      0.9034      5.658      3.562      1.355      5.175 K
ama
6 14.38      14.21      0.8951      5.386      3.312      2.462      4.956 K
ama
> set.seed(123) # (optional) for reproducibility
> seeds_sample <- seeds_original[sample(1:nrow(seeds_original), size = nrow(seeds_original), repl

```

```

ace = TRUE), ]
> dim(seeds_sample)
[1] 210 8
> # Create training and testing sets
> train_index_seeds <- sample(1:nrow(seeds_sample), size = 0.7 * nrow(seeds_sample))
>
> # Training set
> seeds_train <- seeds_sample[train_index_seeds, ]
>
> # Testing set
> seeds_test <- seeds_sample[-train_index_seeds, ]
>
> # Check sizes
> dim(seeds_train) # Should be ~147 rows
[1] 147 8
> dim(seeds_test) # Should be ~63 rows
[1] 63 8
> head(seeds_sample)
  Area Perimeter Compactness LengthKernel WidthKernel AsymmetryCoefficient LengthKernelGroove
Class
159 11.75    13.52    0.8082      5.444      2.678      4.378      5.310
Canadian
207 11.23    12.88    0.8511      5.140      2.795      4.325      5.003
Canadian
179 11.48    13.05    0.8473      5.180      2.758      5.876      5.002
Canadian
14  13.78    14.06    0.8759      5.479      3.156      3.136      4.872
Kama
195 12.11    13.27    0.8639      5.236      2.975      4.132      5.012
Canadian
170 11.24    13.00    0.8359      5.090      2.715      3.521      5.088
Canadian
> # Train the Decision Tree model
> seeds_dt_model <- rpart(Class ~ ., data = seeds_train, method = "class")
>
> # View a quick summary
> summary(seeds_dt_model)
Call:
rpart(formula = Class ~ ., data = seeds_train, method = "class")
n= 147

      CP nsplit  rel error    xerror    xstd
1 0.50561798    0 1.00000000 1.0000000 0.06658254
2 0.38202247    1 0.49438202 0.4943820 0.06238731
3 0.04494382    2 0.11235955 0.1573034 0.03998906
4 0.01000000    3 0.06741573 0.1460674 0.03867901

Variable importance
      Area          Perimeter      WidthKernel      LengthKernel      LengthKerne
lGroove
      19              18              18              17
      13
      Compactness AsymmetryCoefficient
      8              6

Node number 1: 147 observations, complexity param=0.505618
predicted class=Rosa expected loss=0.6054422 P(node) =1
class counts: 40 49 58
probabilities: 0.272 0.333 0.395
left son=2 (91 obs) right son=3 (56 obs)
Primary splits:
  LengthKernelGroove < 5.6185 to the left, improve=46.93367, (0 missing)
  LengthKernel < 5.7695 to the left, improve=44.06224, (0 missing)
  Area < 15.32 to the left, improve=43.43004, (0 missing)
  WidthKernel < 3.393 to the left, improve=43.43004, (0 missing)
  Perimeter < 14.87 to the left, improve=41.97138, (0 missing)
Surrogate splits:
  LengthKernel < 5.7695 to the left, agree=0.986, adj=0.964, (0 split)
  Area < 16.21 to the left, agree=0.973, adj=0.929, (0 split)
  Perimeter < 15.17 to the left, agree=0.973, adj=0.929, (0 split)

```

```

WidthKernel < 3.393 to the left, agree=0.939, adj=0.839, (0 split)
Compactness < 0.8737 to the left, agree=0.680, adj=0.161, (0 split)

```

Node number 2: 91 observations, complexity param=0.3820225

```

predicted class=Kama expected loss=0.4725275 P(node) =0.6190476

```

```

class counts: 40 48 3

```

```

probabilities: 0.440 0.527 0.033

```

```

left son=4 (46 obs) right son=5 (45 obs)

```

Primary splits:

```

Area < 13.395 to the left, improve=31.96522, (0 missing)

```

```

WidthKernel < 3.013 to the left, improve=31.60675, (0 missing)

```

```

Perimeter < 13.985 to the left, improve=27.73469, (0 missing)

```

```

Compactness < 0.86525 to the left, improve=22.33333, (0 missing)

```

```

LengthKernel < 5.478 to the left, improve=20.87857, (0 missing)

```

Surrogate splits:

```

Perimeter < 13.75 to the left, agree=0.967, adj=0.933, (0 split)

```

```

WidthKernel < 3.102 to the left, agree=0.945, adj=0.889, (0 split)

```

```

LengthKernel < 5.4745 to the left, agree=0.879, adj=0.756, (0 split)

```

```

Compactness < 0.869 to the left, agree=0.835, adj=0.667, (0 split)

```

```

AsymmetryCoefficient < 3.592 to the right, agree=0.747, adj=0.489, (0 split)

```

Node number 3: 56 observations

```

predicted class=Rosa expected loss=0.01785714 P(node) =0.3809524

```

```

class counts: 0 1 55

```

```

probabilities: 0.000 0.018 0.982

```

Node number 4: 46 observations, complexity param=0.04494382

```

predicted class=Canadian expected loss=0.1304348 P(node) =0.3129252

```

```

class counts: 40 6 0

```

```

probabilities: 0.870 0.130 0.000

```

```

left son=8 (38 obs) right son=9 (8 obs)

```

Primary splits:

```

AsymmetryCoefficient < 3.263 to the right, improve=7.434783, (0 missing)

```

```

LengthKernelGroove < 4.9805 to the right, improve=6.434783, (0 missing)

```

```

Compactness < 0.8656 to the left, improve=3.490338, (0 missing)

```

```

WidthKernel < 3.013 to the left, improve=2.645309, (0 missing)

```

```

LengthKernel < 5.167 to the right, improve=2.034783, (0 missing)

```

Surrogate splits:

```

LengthKernelGroove < 4.7885 to the right, agree=0.957, adj=0.750, (0 split)

```

```

Compactness < 0.8656 to the left, agree=0.913, adj=0.500, (0 split)

```

```

WidthKernel < 3.113 to the left, agree=0.891, adj=0.375, (0 split)

```

Node number 5: 45 observations

```

predicted class=Kama expected loss=0.06666667 P(node) =0.3061224

```

```

class counts: 0 42 3

```

```

probabilities: 0.000 0.933 0.067

```

Node number 8: 38 observations

```

predicted class=Canadian expected loss=0 P(node) =0.2585034

```

```

class counts: 38 0 0

```

```

probabilities: 1.000 0.000 0.000

```

Node number 9: 8 observations

```

predicted class=Kama expected loss=0.25 P(node) =0.05442177

```

```

class counts: 2 6 0

```

```

probabilities: 0.250 0.750 0.000

```

```
> # Predict classes for training data
```

```
> seeds_train_pred <- predict(seeds_dt_model, seeds_train, type = "class")
```

```
>
```

```
> # Predict classes for testing data
```

```
> seeds_test_pred <- predict(seeds_dt_model, seeds_test, type = "class")
```

```
>
```

```
> # Build confusion matrices
```

```
> seeds_train_conf_matrix <- table(Predicted = seeds_train_pred, Actual = seeds_train$Class)
```

```
> seeds_test_conf_matrix <- table(Predicted = seeds_test_pred, Actual = seeds_test$Class)
```

```
>
```

```
> # View matrices
```

```
> seeds_train_conf_matrix
```

```
  Actual
```

```

Predicted   Canadian Kama Rosa
Canadian    38     0    0
Kama        2    48    3
Rosa        0     1   55
> seeds_test_conf_matrix
      Actual
Predicted Canadian Kama Rosa
Canadian    25     3    0
Kama        3    17    1
Rosa        0     0   14
> seeds_train_accuracy <- sum(diag(seeds_train_conf_matrix)) / sum(seeds_train_conf_matrix)
> seeds_train_accuracy
[1] 0.9591837
> seeds_test_accuracy <- sum(diag(seeds_test_conf_matrix)) / sum(seeds_test_conf_matrix)
> seeds_test_accuracy
[1] 0.8888889
> # Function to calculate precision, recall (sensitivity), and F1
> calc_metrics <- function(conf_matrix) {
+   precision <- diag(conf_matrix) / rowSums(conf_matrix)
+   recall <- diag(conf_matrix) / colSums(conf_matrix)
+   f1 <- 2 * precision * recall / (precision + recall)
+   data.frame(Precision = precision, Recall = recall, F1 = f1)
+ }
>
> # Sensitivity (Recall) for Training set
> seeds_train_metrics <- calc_metrics(seeds_train_conf_matrix)
> seeds_train_metrics
      Precision    Recall      F1
Canadian 1.0000000 0.9500000 0.9743590
Kama     0.9056604 0.9795918 0.9411765
Rosa     0.9821429 0.9482759 0.9649123
>
> # Sensitivity (Recall) for Testing set
> seeds_test_metrics <- calc_metrics(seeds_test_conf_matrix)
> seeds_test_metrics
      Precision    Recall      F1
Canadian 0.8928571 0.8928571 0.8928571
Kama     0.8095238 0.8500000 0.8292683
Rosa     1.0000000 0.9333333 0.9655172
> # Function to calculate specificity
> calc_specificity <- function(conf_matrix) {
+   spec <- numeric(nrow(conf_matrix))
+   for (i in 1:nrow(conf_matrix)) {
+     true_negatives <- sum(conf_matrix[-i, -i])
+     false_positives <- sum(conf_matrix[i, -i])
+     spec[i] <- true_negatives / (true_negatives + false_positives)
+   }
+   return(spec)
+ }
>
> # Specificity for Training set
> seeds_train_specificity <- calc_specificity(seeds_train_conf_matrix)
> seeds_train_specificity
[1] 1.0000000 0.9489796 0.9887640
>
> # Specificity for Testing set
> seeds_test_specificity <- calc_specificity(seeds_test_conf_matrix)
> seeds_test_specificity
[1] 0.9142857 0.9069767 1.0000000
> library(caret)
> # Create a vector to store accuracies
> seeds_accuracies <- c()
>
> # List of k values
> k_values <- c(5, 10, 15, 20)
>
> # Loop through each k
> for (k in k_values) {
+
+   # Set up k-fold cross validation

```

```
+ cv_control_seeds <- trainControl(method = "cv", number = k)
+
+ # Train the Decision Tree model using caret's train() function
+ cv_model_seeds <- train(Class ~ .,
+                          data = seeds_sample,
+                          method = "rpart",
+                          trControl = cv_control_seeds)
+
+ # Store the mean accuracy
+ seeds_accuracies <- c(seeds_accuracies, max(cv_model_seeds$results$Accuracy))
+ }
>
> # View the results
> data.frame(K = k_values, Accuracy = seeds_accuracies)
  K Accuracy
1  5 0.9192027
2 10 0.9290043
3 15 0.9166789
4 20 0.9172727
> # Plot Accuracy vs K
> plot(k_values, seeds_accuracies, type = "b", pch = 19,
+       xlab = "K (Number of Folds)",
+       ylab = "Accuracy",
+       main = "Seeds Dataset - Accuracy vs K for Decision Tree",
+       ylim = c(0, 1))
>
```